



Self-Checkpointing

Tim Cartwright

*OSG Deputy Executive Director and User School Director
University of Wisconsin–Madison*



The Challenge

- Suppose your job will run for a long time
 - Reminder: Look at the “Ideal Jobs” table
 - But let’s say more than about 8 hours
- Likely removed from the Execution Point before done: HTCondor will restart job somewhere else
- ***But!*** It starts over and loses all progress (*badput*)



Some Solutions

- **Ideal solution:** Break up job into shorter pieces
 - Try to get back into that “Ideal Jobs” column
- But this does not always work; for example, when one iteration depends on the previous one
- Another solution: **Self-checkpointing**

The executable periodically saves its progress to disk – a *self-made checkpoint* – so that it can resume from that point if interrupted later, losing minimal progress



Requirements

- Your executable can self-checkpoint and resume progress from checkpoint file(s) upon restart
 - If you have the source code, you can probably do this
 - If not, the code must have the feature already
 - A wrapper script *may* be able to help, but seems tricky
- Using HTCondor $\geq 9.0.6$ is good; $\geq 9.10.0$ is best
 - CHTC and OSPool are both $\geq 9.10.0$
- Job universe: vanilla (default) or Docker (container)



HTCondor Has 2 Ways to Checkpoint

- **Exit-driven self-checkpointing**
 - Since HTCondor $\geq 8.9.7$
 - *Waaaay* better for most use cases, esp. in OSG
 - What is shown here
- Eviction-driven self-checkpointing
 - Not even worth talking about for OSG!
 - Documented in the HTCondor Manual
 - But don't use it 😊



Technical Details



HTCondor Submit File Changes

- Tell HTCondor what special exit code your software will use when checkpointing (85 is suggested):

```
checkpoint_exit_code = 85
```

- Tell HTCondor what files (on the Execution Point) to save (on the Access Point) and restore *when moved to a new Execution Point* — list files and directories, include output file(s) if cumulative:

```
transfer_checkpoint_files = foo.txt, ...
```




Example Submit File

```
executable = my_software
transfer_input_files      = my_input.txt
transfer_checkpoint_files = my_output.txt, temp_dir, temp_file.txt
transfer_output_files     = my_output.txt
```

```
request_cpus    = 1
request_memory  = 1GB
request_disk    = 1GB
```

```
log      = example.log
output   = example.out
error    = example.err
```

```
checkpoint_exit_code = 85
```

```
queue
```




Notes About Checkpointed Files

- If you omit **transfer_checkpoint_files**, HTCondor uses **transfer_output_files** (or its defaults)
- Consider Access Point storage needs; can estimate as:
number of running jobs × total size of checkpoint files
(OSPool uses your **/home** quota; elsewhere: ask admin)
- So, save only what you need! Because it identifies exact files, it can help to use **transfer_checkpoint_files**



Executable (Code) Changes

- Executable may run many times before finishing; external process (HTCondor) reruns it until *done*
- Periodically write state to file(s), then immediately exit with **transfer_checkpoint_files** (85)
- Any other exit code indicates *done* (good or error)
- At start-up, executable must check for presence of checkpoint file(s); **if absent**, start at beginning, but **if present**, read file(s) and resume from that point



Self-Checkpoint Frequency

- Balance overhead versus (risk of) lost computing
 - Writing to disk can be slow and restarts take time
 - Test early! Collect metrics (checkpoint & restart times)
- Look for natural checkpoint times
 - Generally, when there is the least data to write
 - Often between outermost iterations
 - Could use iteration count, time, ...
- As a starting point, checkpoint every 1–2 hours



Debugging Tips

- For testing, you can force HTCondor to stop your job and run again (new sandbox, maybe new EP):

condor_vacate_job *JobID*

- If HTCondor has transferred checkpoint files back to the Access Point, you can get a copy with:

condor_evicted_files get *JobID*



Step-by-Step Example



Example Step 1: Before Submit

Submit Directory

```
my_software  
my_input.txt  
my_submit.sub
```

```
executable = my_software  
transfer_input_files = my_input.txt  
transfer_checkpoint_files = my_output.txt, temp_dir,  
                           temp_file.txt  
transfer_output_files = my_output.txt  
  
request_cpus      = 1  
request_memory    = 1GB  
request_disk      = 1GB  
  
log               = zzz.log  
output            = zzz.out  
error             = zzz.err  
  
checkpoint_exit_code = 85  
  
queue
```



Example Step 2: Just Before Execute


Submit Directory

my_software
my_input.txt
my_submit.sub
zzz.log

Spool Directory

Execute Directory

my_input.txt
my_software





Example Step 3: After 1 Minute

Submit Directory

```
my_software  
my_input.txt  
my_submit.sub  
zzz.log
```

Spool Directory

Execute Directory

```
my_input.txt  
my_output.txt  
my_software  
_condor_stderr  
_condor_stdout  
temp-dir/1.txt  
temp-dir/2.txt  
temp-file.txt  
trash.txt
```



Example Step 4: After 1 Hour – `exit(85)`

Submit Directory

```
my_software  
my_input.txt  
my_submit.sub  
zzz.log
```

Spool Directory

Execute Directory

```
my_input.txt  
my_output.txt  
my_software  
_condor_stderr  
_condor_stdout  
temp-dir/42.txt  
temp-dir/43.txt  
temp-file.txt  
trash.txt
```



Example Step 5: Checkpoint Complete

```
transfer_checkpoint_files = my_output.txt, temp-dir, temp-file.txt
```

Submit Directory


```
my_software  
my_input.txt  
my_submit.sub  
zzz.log
```

Spool Directory

```
my_output.txt  
_condor_stderr  
_condor_stdout  
temp-dir/42.txt  
temp-dir/43.txt  
temp-file.txt
```

Execute Directory

```
my_input.txt  
my_output.txt  
my_software  
_condor_stderr  
_condor_stdout  
temp-dir/42.txt  
temp-dir/43.txt  
temp-file.txt  
trash.txt
```



Job execute directory is not
changed before restart.



Example Step 6: 10 Min. Later – Eviction!

Submit Directory

```
my_software  
my_input.txt  
my_submit.sub  
zzz.log
```

Spool Directory

```
my_output.txt  
temp-dir/42.txt  
temp-dir/43.txt  
temp-file.txt
```

Execute Directory

```
my_input.txt  
my_output.txt  
my_software  
_condor_stderr  
_condor_stdout  
temp-dir/51.txt  
temp-dir/52.txt  
temp-file.txt  
trash.txt
```

Lose changes since last checkpoint



Example Step 7: Restart on New Execute

Submit Directory


```
my_software  
my_input.txt  
my_submit.sub  
zzz.log
```

Spool Directory

```
my_output.txt  
_condor_stderr  
_condor_stdout  
temp-dir/42.txt  
temp-dir/43.txt  
temp-file.txt
```

(New) Execute Directory

```
my_input.txt  
my_output.txt  
my_software  
_condor_stderr  
_condor_stdout  
temp-dir/42.txt  
temp-dir/43.txt  
temp-file.txt
```





Example Step 8: Job Completes Normally


```
transfer_output_files = my_output.txt
```

Submit Directory

my_software
my_input.txt
my_output.txt
my_submit.sub
zzz.err
zzz.log
zzz.out

(New) Execute Directory

my_input.txt
my_output.txt
my_software
_condor_stderr
_condor_stdout
temp-dir/98.txt
temp-dir/99.txt
temp-file.txt
trash.txt





Notes & Acknowledgements

- Official documentation:
 - <https://htcondor.readthedocs.io/en/latest/users-manual/self-checkpointing-applications.html>
 - Includes full working example (Python + submit)
 - The exercise is derived from that example
- Many thanks to Todd Miller, Christina Koch, and Jason Patton for their help!
- This work was supported by NSF grants MPS-1148698, OAC-1836650, and OAC-2030508